

Resolution Semantics for Binary Partial Algebras

A Formalized Free Completion, Intrinsic Finite-Complement Separation, and Observational Completion

Adrian Puha

August 2026

Abstract

We formalize a normalized free compatible completion of binary partial algebras. Old-defined applications reduce to old values; old-undefined applications become structured finite-tree Answers. The Answer algebra has the expected universal property and is the expression algebra modulo equality after resolution.

Our additional construction is an intrinsic finite-complement separating semantics. Every distinct pair is distinguished by a compatible total extension fixing the old carrier pointwise whose complement injects into $\text{Fin}(n + 1)$. At each budget this is equivalent to the canonical presentation

$$D.\text{Carrier} \sqcup \text{Fin}(n) \sqcup \{\star\}.$$

The fresh-tag count is at most the combined constructor size; this is a construction bound, not an optimality claim. When the old carrier is finitely coded, the same witnesses are finite in the ordinary sense, so the generated Answer algebra is residually finite with an explicit total-state bound.

Finite-tag agreement induces a separated filtration whose Cauchy completion supports canonical operations and has the same universal equational theory. For a finitely coded old carrier, the embedding is surjective exactly when the old evaluator is total. Without any finiteness assumption, an undefined old application plus a context $x \mapsto u(x, e)$ fixing every old element yields a factorial Cauchy orbit with no generated limit. Every finite stage is strictly coarser than equality, and least separating budgets are unbounded. Natural and integer arithmetic satisfy this criterion using $0/0$ and $x \mapsto x + 0$, so both completion embeddings are proper. Division by zero remains a structured non-old Answer, not an ordinary number. Lean 4.33.0 checks the development. The exact priority of this combined package remains under specialist review.

Contents

1	Introduction	2
2	Generated Answers as a Free Compatible Completion	4
2.1	Expressions and Answers	4
2.2	Universal property	5
2.3	Quotient presentation	5

3	Intrinsic Finite-Complement Separation	6
3.1	Pair-dependent separator	7
4	The Observational Filtration and Completion	8
4.1	Relative universal property	11
5	Exact Equational Conservativity	11
6	Arithmetic Examples	12
6.1	Natural numbers	12
6.2	Integers	13
7	Relation to Prior Work	13
7.1	Free completion and partial algebras	13
7.2	Operation trees and effects	13
7.3	Tree metrics and completion	13
7.4	Finite complements and profinite comparison	14
7.5	Totalized division	15
7.6	Candidate contribution	15
8	Formalization and Trust Boundary	16
9	Limitations and Open Questions	16
10	Conclusion	17

1 Introduction

Let D be a one-sorted binary partial algebra with carrier A , operation symbols Σ , and evaluator

$$\text{eval}_D : \Sigma \longrightarrow A \longrightarrow A \longrightarrow \text{Option}(A).$$

The value $\text{some}(c)$ records that the old application is defined and equals c ; none records old undefinedness.

There is extensive prior work on completing or encoding partial algebras. Climent Vidal and Cosme Llópez prove that the inclusion of total algebras into partial algebras has a left adjoint, the absolutely free completion functor, and develop the functorial Schmidt construction [VL23]. Diaconescu gives a Horn-style encoding of partial algebras as total algebras [Dia09]. Completion questions also occur in partial combinatory algebra and term-rewriting models [BKdV96, BKdV99, Nes11]. Accordingly, this paper does not claim that the existence of a free or total extension is new.

Our starting presentation is operationally simple. Resolve an expression bottom-up. If a node has two old operands and the old evaluator returns a value, replace the node by that old value. Otherwise retain the node, its operation symbol, and its resolved children. The resulting finite trees are called *Answers*. This gives a normalized representative for a free compatible total extension.

The paper asks two narrower questions.

- (Q1) Can equality of generated Answers be characterized by total models that do not themselves contain Answer trees, while preserving every old-defined application?
- (Q2) Can those finite observations support a completion on which the operations remain canonical and the universal equational theory is unchanged?

The answer to both is yes. The paper has two principal results.

Theorem 1.1 (Structured free completion and finite separation). *For every one-sorted binary partial algebra D :*

- (a) *the generated Answer algebra $R(D)$ is initial among compatible total algebras equipped with a map from the old carrier;*
- (b) *$R(D)$ is represented by the quotient of the absolutely free expression algebra by the kernel of the resolution evaluator;*
- (c) *every distinct $x, y \in R(D)$ is separated by a compatible total extension whose complement over A has at most $|x| + |y| + 1$ points; equivalently, a canonical finite-tag model needs at most $|x| + |y|$ fresh tags and one overflow state. Consequently, when A is finite, $R(D)$ is residually finite.*

Theorem 1.2 (Observational completion and equation preservation). *Finite-tag agreement defines a separated descending filtration on $R(D)$. Its Cauchy completion $\widehat{R(D)}$ is complete, contains an injective finite-stage dense image of $R(D)$, and carries the unique jointly nonexpansive extensions of the primitive operations. Relative to complete compatible filtered targets, the completed algebra has the expected universal extension property. If the old carrier is finite, the canonical embedding is surjective exactly when the old evaluator is total; equivalently, the completion is proper exactly when some old application is undefined. Without assuming the old carrier finite, suppose some old application is undefined and there are $u \in \Sigma$ and $e \in A$ with $\text{eval}_D(u, a, e) = \text{some}(a)$ for every $a \in A$. Then every finite observation stage is strictly coarser than equality, an explicit factorial orbit is Cauchy with no generated limit, its separation ranks are unbounded, and the completion embedding is not surjective. Moreover, for every pair of constant-bearing algebraic terms s, t ,*

$$\widehat{R(D)} \models s = t \iff R(D) \models s = t.$$

The first two clauses of [theorem 1.1](#) place the generated construction within established free-completion theory. The more specific ingredients are the intrinsic finite-complement observer class over a pointwise-preserved base, its pairwise separator, and the observational structure it induces. The linear estimate is retained as a general construction bound rather than a state-complexity advance. Over a finitely coded base it also gives ordinary residual finiteness and an explicit bound on the entire separating target. The finite-base properness criterion explains exactly when the Cauchy construction adds genuinely new points over a finite old carrier; the old-fixing-context criterion shows that properness is not confined to finite bases. The two hypotheses are incomparable: the one-point algebra is finite and

partial but has no old-fixing context, whereas natural arithmetic has an old-fixing context but no finite code for its carrier. Natural and integer arithmetic are concrete corollaries. We have not found the exact combined package in the literature reviewed so far, but this is not a claim of historical priority.

2 Generated Answers as a Free Compatible Completion

2.1 Expressions and Answers

Finite input expressions and raw Answers are defined by separate grammars:

$$e ::= \text{val}(a) \mid \text{app}(f, e_1, e_2),$$

$$r ::= \text{old}(a) \mid \text{susp}(f, r_1, r_2).$$

The separation is intentional. Expressions are input syntax. Answers are normal forms produced by the partial evaluator together with the suspension rule.

Definition 2.1 (One-step resolution). For $f \in \Sigma$ and raw Answers x, y , let

$$\text{liftOp}_D(f, x, y) = \begin{cases} \text{old}(c), & x = \text{old}(a), y = \text{old}(b), \text{eval}_D(f, a, b) = \text{some}(c), \\ \text{susp}(f, x, y), & \text{otherwise.} \end{cases}$$

The resolution function Res_D is the structural recursion determined by

$$\text{Res}_D(\text{val}(a)) = \text{old}(a),$$

$$\text{Res}_D(\text{app}(f, e_1, e_2)) = \text{liftOp}_D(f, \text{Res}_D(e_1), \text{Res}_D(e_2)).$$

The generated carrier is the image of this evaluator:

$$R(D) = \{r \mid \exists e, \text{Res}_D(e) = r\}.$$

The old embedding is $\eta_0(a) = \text{old}(a)$. Generated operations are

$$f_{R(D)}(x, y) = \text{liftOp}_D(f, x, y).$$

Generation is closed under these operations because expression witnesses can be joined by one application node.

Proposition 2.2 (Conservativity on defined applications). *If $\text{eval}_D(f, a, b) = \text{some}(c)$, then*

$$f_{R(D)}(\eta_0(a), \eta_0(b)) = \eta_0(c).$$

The map η_0 is injective.

The proposition is immediate from [definition 2.1](#). Its role is not depth but control: no old-defined calculation is reinterpreted.

2.2 Universal property

A *compatible total algebra over D* consists of a total Σ -algebra T , a map $i : A \rightarrow T$, and the preservation law

$$\text{eval}_D(f, a, b) = \text{some}(c) \implies f_T(i(a), i(b)) = i(c).$$

The map i is not required to be injective. This is the target notion needed for the ordinary free-completion universal property.

Every raw Answer has an interpretation in T :

$$\Phi_T(\text{old}(a)) = i(a), \quad \Phi_T(\text{susp}(f, x, y)) = f_T(\Phi_T(x), \Phi_T(y)).$$

The preservation law gives

$$\Phi_T(\text{liftOp}_D(f, x, y)) = f_T(\Phi_T(x), \Phi_T(y)).$$

The only non-structural case is the old-defined branch, where the equation is exactly the compatibility condition.

Theorem 2.3 (Free compatible completion). *For every compatible total algebra T over D , there is a unique homomorphism*

$$\Phi_T : R(D) \longrightarrow T$$

such that $\Phi_T \circ \eta_0 = i$.

Proof. Existence is the fold above. It preserves generated operations by the one-step equation. For uniqueness, choose an expression e generating a given Answer. Induction on e forces every compatible homomorphism to agree with the fold on $\text{Res}_D(e)$. Since every generated Answer has such a witness, the two maps agree everywhere. $\square$

The Lean-facing theorem is `ResolutionSemantics.freeCompletionUniversal`. This result confirms that the generated construction belongs to free-completion theory rather than constituting a new category-theoretic notion.

2.3 Quotient presentation

Define an equivalence on input expressions by

$$e \sim_D e' \iff \text{Res}_D(e) = \text{Res}_D(e').$$

Then resolution descends to a map

$$\overline{\text{Res}}_D : \text{Expr}(D) / \sim_D \longrightarrow R(D).$$

Proposition 2.4 (Kernel quotient presentation). *The map $\overline{\text{Res}}_D$ is injective and surjective. Thus generated Answers are a normalized presentation of the expression algebra modulo the resolution kernel.*

The two inverse laws are formalized by `ResolutionSemantics.expressionKernelQuotientInjective` and `ResolutionSemantics.expressionKernelQuotientSurjective`.

Remark 2.5 (Relation to the Schmidt construction). The Schmidt construction and the left adjoint free-completion functor are direct prior art [VL23]. Our formal theorem proves the corresponding universal property for the concrete one-sorted binary presentation used here. A fully explicit categorical isomorphism between the Lean object and every formulation of the many-sorted Schmidt construction is not needed for the results below and is not claimed. The correct positioning is that $R(D)$ is a normalized free compatible completion.

3 Intrinsic Finite-Complement Separation

The free-completion theorem alone does not show that distinctions between suspended Answers are visible in independently specified models. The primitive observer condition should not depend on a particular sum encoding.

Let T be a compatible total extension with injective old embedding $i : A \hookrightarrow T$. Define its genuine outside part by

$$\text{Out}_D(T) = \{z \in T \mid \forall a \in A, z \neq i(a)\}.$$

Definition 3.1 (Intrinsic finite-complement observer). An observer has budget n when there is an injection

$$\text{Out}_D(T) \hookrightarrow \text{Fin}(n + 1).$$

Thus only the complement over the pointwise-preserved base is finite; the old carrier itself may be infinite. The budget is an upper bound, so no ordering, chosen tags, or distinguished overflow point is part of this definition.

Every compatible extension has the carrier decomposition

$$T \cong A \sqcup \text{Out}_D(T).$$

For the concrete observer used in the construction below, the outside part is canonically $\text{Fin}(n) \sqcup \{\star\}$.

Definition 3.2 (Finite-tag model). For $n \in \mathbb{N}$, a finite-tag model over D has carrier

$$A \sqcup \text{Fin}(n) \sqcup \{\star\}$$

and total operations preserving every old-defined application:

$$\text{eval}_D(f, a, b) = \text{some}(c) \implies f_T(\text{old}(a), \text{old}(b)) = \text{old}(c).$$

No further equations are imposed.

The formal representation theorem proves that intrinsic separation at budget n is equivalent to finite-tag separation at the same budget. In the forward direction, an intrinsic complement code yields a base-fixing injection into the canonical carrier. Unused tag states are allowed and are sent to overflow when the operations are transported. This is an injection presentation, not a claimed bijection with all canonical states.

Write $x \equiv_n y$ when every finite-tag model with n tags interprets x and y equally. Full finite observational agreement means

$$x \equiv_{\text{fin}} y \iff \forall n, x \equiv_n y.$$

3.1 Pair-dependent separator

Let $x \neq y$ be generated Answers. Form a finite list $S(x, y)$ containing all subterms of the two raw trees. Repetition is harmless; first-occurrence indices provide a finite set of tags. Define an encoding on selected nodes:

$$q(\text{old}(a)) = \text{old}(a), \quad q(s) = \text{tag}(\text{index}(s))$$

for selected suspended nodes s . The encoding is injective on the selected closure.

The model operation first preserves an old-defined application. Otherwise it attempts to decode both inputs, constructs one suspended node, and re-encodes that node if it belongs to the selected closure; all unrelated cases go to $\star$. There is no branch conflict. A generated normal Answer cannot contain a suspended node with two old children whose old application is defined, because resolution would already have reduced it.

Induction on selected subterms gives

$$\Phi_T(s) = q(s).$$

In particular, $\Phi_T(x) = q(x)$ and $\Phi_T(y) = q(y)$, so injectivity of q on $S(x, y)$ separates the roots.

Theorem 3.3 (Intrinsic finite-complement separation). *For distinct $x, y \in R(D)$ there exist n and a compatible total extension T whose outside part injects into $\text{Fin}(n + 1)$ and such that*

$$\Phi_T(x) \neq \Phi_T(y).$$

One may choose

$$n = |x| + |y|,$$

where $|-|$ is raw constructor count. Consequently,

$$x \equiv_{\text{fin}} y \iff x = y.$$

The intrinsic public Lean declarations are `ResolutionSemantics.qualitativeFiniteComplementSeparating_theorem`, `ResolutionSemantics.intrinsicSeparationIffFiniteTag`, and `ResolutionSemantics.intrinsicExtensionHasCanonicalPresentation`. The canonical finite-tag consequences remain exposed as `ResolutionSemantics.finiteExternalSeparation`, `ResolutionSemantics.finiteObservationComplete`, and `ResolutionSemantics.finiteSeparationBound`.

Remark 3.4 (Interpretation of the result). The old carrier A may be infinite. The theorem is therefore not finite residuality of the whole carrier. It says intrinsically that a chosen pair requires only finitely many *new* semantic states beyond A ; the finite-tag sum is a derived normal form. The displayed linear estimate is a baseline construction bound, not an optimality theorem. We avoid using “full abstraction” as a novelty label; the mathematical content is the explicit model class and separator.

More precisely, let $\mathcal{C}_D$ be the class of compatible total extensions that fix A pointwise and have finite complement over A . The theorem says that $R(D)$ is residually $\mathcal{C}_D$: every unequal pair is separated by its canonical homomorphism into some member of $\mathcal{C}_D$. This is an exact classification of the statement as a relative residual property, not a claim that the terminology or general residual paradigm is new.

Corollary 3.5 (Finite-base residual finiteness). *Suppose the old carrier admits an injective code into $\text{Fin}(m)$. Then $R(D)$ is residually finite in the ordinary homomorphic sense. More precisely, distinct x, y are separated by a finite compatible algebra with an injective code into*

$$\text{Fin}(m + |x| + |y| + 1).$$

The formal statements are `ResolutionSemantics.ResidualComparison.finiteBaseCompatibleSeparationBound` and `ResolutionSemantics.ResidualComparison.finiteBaseGeneratedResiduallyFinite`. The target produced by the proof still embeds the old carrier injectively, so the formal result is slightly stronger than ordinary residual finiteness, whose separating homomorphisms need not be faithful on the chosen generators.

Proposition 3.6 (Closure of the old image). *Let $\iota_D : A \hookrightarrow R(D)$ be the canonical old embedding. Its image is closed under every total operation of $R(D)$ if and only if the old evaluator is total.*

Indeed, a defined old application reduces to an old value. If $\text{eval}_D(f, a, b) = \text{none}$, however, then

$$f^{R(D)}(\iota_D(a), \iota_D(b)) = \text{susp}(f, \text{old}(a), \text{old}(b)),$$

which is not in the old image. Lean checks both the pointwise escape statement `ResolutionSemantics.ResidualComparison.undefinedApplicationEscapesOldImage` and the equivalence `ResolutionSemantics.ResidualComparison.generatedOldClosedIffEvaluatorTotal`.

Residual finiteness and recognizability of free algebras are established subjects. The present targets are generally not finite: the base carrier A may be infinite and is retained pointwise, while only the added tags and $\star$ are counted. The closest computational analogy is the separating- words problem. Sublinear worst-case bounds were established by Goralcik and Koubek, Robson, and Chase [GK86, Rob89, Cha20]. Along a unary comb, a Resolution observer likewise iterates one transition map from a fixed starting state. This analogy confirms that our linear estimate is a baseline, but it does not transfer word bounds automatically: arbitrary Answer trees branch, and every observer must extend the pointwise-preserved partial algebra. The resulting relative separation complexity is open. Determining whether this exact observer class already has a standard name remains part of our specialist literature audit.

4 The Observational Filtration and Completion

Each $\equiv_n$ is an equivalence relation. The formal model family yields a descending filtration: agreement at a larger observational budget implies agreement at every smaller budget. By [theorem 3.3](#),

$$\bigcap_{n \in \mathbb{N}} \equiv_n = .$$

Thus the filtration is separated.

A symbolic non-Archimedean distance can be read from the least separating stage. Equal points have distance zero; a distinct pair receives the level of its least finite separator. The strong triangle law follows from transitivity of every stage relation. This is standard filtered-space machinery, analogous to established metric treatments of trees [AN80, Cou83].

A sequence (x_k) is Cauchy when, for every stage n , its tail is pairwise $\equiv_n$ -equivalent. Two Cauchy sequences are asymptotically equivalent when they eventually agree at every fixed stage. Let

$$\widehat{R(D)} = \{\text{Cauchy sequences in } R(D)\} / \sim.$$

The constant-sequence map is injective by separatedness and stage-dense by the usual late-term argument. A diagonal construction proves completeness.

Theorem 4.1 (Finite-base properness criterion). *Suppose the old carrier A admits an injection into $\text{Fin}(m)$ for some m . Then the canonical embedding*

$$R(D) \longrightarrow \widehat{R(D)}$$

is surjective if and only if the old evaluator is total. Equivalently, this embedding is not surjective if and only if there exist $f \in \Sigma$ and $a, z \in A$ with $\text{eval}_D(f, a, z) = \text{none}$.

Proof. A stage- n observer has at most $m + n + 1$ states: the finitely coded old carrier, n fresh tags, and overflow. If $\text{eval}_D(f, a, z)$ is undefined, let c_k be the left comb obtained by iterating this application. Along the comb family every observer acts by the orbit of one self-map. Hence $(c_k!)_{k \in \mathbb{N}}$ is Cauchy: sufficiently late factorial indices have difference divisible by every possible orbit period.

For non-convergence, fix a proposed finite Answer x . Choose the preserving finite-tag observer constructed from the finite subterm closure of x together with a comb prefix longer than x . It encodes x and the prefix by distinct non-overflow states. The next comb transition matches no selected normalized node, so it enters overflow; no selected transition has overflow as an encoded child, making overflow absorbing along the rest of the comb. Thus this single finite stage separates x from every sufficiently late factorial comb. No generated Answer is a limit.

Conversely, if the old evaluator is total, structural induction shows that every generated Answer is old. Cauchy agreement at a fixed separating stage then makes every Cauchy sequence eventually constant, so the generated space is complete. Completeness is equivalent to surjectivity of the canonical completion embedding. The checked public statements are `ResolutionSemantics.FiniteBaseCompletion.completeIffTotal` and `ResolutionSemantics.FiniteBaseCompletion.embeddingNotSurjectiveIffExistsUndefined`. $\square$

The one-point always-undefined algebra is the smallest corollary and still provides an explicit represented completion point through `ResolutionSemantics.OnePointCompletion.addsPoint`. The theorem does not by itself decide infinite-base examples. The following independent criterion does.

Theorem 4.2 (Old-fixing-context properness). *Let D be any binary partial algebra; no finiteness assumption is made on its old carrier. Suppose*

$$\text{eval}_D(g, a, b) = \text{none} \quad \text{and} \quad \text{eval}_D(u, x, e) = \text{some}(x) \quad (x \in A)$$

for some $g, u \in \Sigma$ and $a, b, e \in A$. Define

$$c_0 = \text{susp}(g, \text{old}(a), \text{old}(b)), \quad c_{k+1} = \text{susp}(u, c_k, \text{old}(e)).$$

Then:

- (a) $c_1 \neq c_2$ but $c_1 \equiv_0 c_2$;
- (b) for every m , the distinct pair $c_{(m+2)!}, c_{(m+3)!}$ agrees at stage m , and its least finite-tag separation rank is greater than m ;
- (c) $(c_{k!})_{k \in \mathbb{N}}$ is Cauchy and has no generated limit.

Consequently the generated filtered space is not complete, the finite-tag separation ranks are unbounded, and the canonical completion embedding is not surjective.

Proof. Let T be a stage- m preserving observer and let $F(s) = u_T(s, \text{old}(e))$. Preservation gives $F(\text{old}(x)) = \text{old}(x)$ for every $x \in A$. For this orbit, collapse the entire old carrier to one settled state while retaining the m fresh tags and overflow. The collapse commutes with F and has $m + 2$ states. Equality in the collapsed orbit lifts to the original orbit: complement states were retained injectively, while an orbit that reaches an old value $\text{old}(x)$ stays at that exact value because $F(\text{old}(x)) = \text{old}(x)$. Finite-orbit periodicity therefore implies that sufficiently late indices whose difference is divisible by $(m + 2)!$ have equal T -values. Applying this simultaneously to factorial indices proves that $(c_{k!})$ is Cauchy and gives the stage- m agreement in clause (b). Constructor count proves the pair distinct. Since a distinct pair agrees at stage m , the exact rank characterization makes its least separating budget greater than m .

At stage zero the complement consists only of overflow. If F sends overflow to an old value, the next application fixes that value; if it sends overflow to overflow, the same table entry fixes it. Thus $F^2 = F$, proving $c_1 \equiv_0 c_2$, while constructor count again proves $c_1 \neq c_2$.

For non-convergence, fix a proposed finite Answer x . The standard candidate-tailored preserving observer encodes x together with a repeated $u(-, e)$ prefix longer than x . The next iterate is absent from the selected finite table and therefore enters overflow; overflow is absorbing along the remaining tail. The candidate remains encoded by a non-overflow state, so one finite stage separates x from every sufficiently late factorial term. Hence no generated Answer is a limit. The checked public statements are `ResolutionSemantics.OldFixingContextCompletion.factorialSequenceCauchy`, `ResolutionSemantics.OldFixingContextCompletion.factorialSequenceNoLimit`, `ResolutionSemantics.OldFixingContextCompletion.separationRankGreaterThanBudget`, and `ResolutionSemantics.OldFixingContextCompletion.embeddingNotSurjective`. $\square$

Natural and integer arithmetic both satisfy the hypotheses with undefined seed 0/0 and old-fixing context $x \mapsto x + 0$. Thus [theorem 4.2](#) proves that both arithmetic completion embeddings are non-surjective and that neither filtration has a finite stage equal to equality. The direct public corollaries include `ResolutionSemantics.NatDivision.oldFixingCriterion`, `ResolutionSemantics.NatDivision.separationRanksUnbounded`, `ResolutionSemantics.IntDivision.completionEmbeddingNotSurjective`, and `ResolutionSemantics.IntDivision.separationRanksUnbounded`.

Proposition 4.3 (Incomparability of the properness hypotheses). *Neither properness hypothesis subsumes the other.*

- (i) *The one-point always-undefined algebra has a one-state code, an undefined old application, and a proper completion, but it admits no old-fixing-context witness.*
- (ii) *Natural arithmetic has the old-fixing witness $x \mapsto x + 0$ and a proper completion, but its old carrier admits no finite code.*

Proof. For the one-point algebra, any old-fixing witness would require its sole operation to satisfy $\text{eval}_D(u, (), ()) = \text{some}(()),$ contradicting the definition $\text{eval}_D \equiv \text{none}.$ Its one-state code, undefined operation, and proper completion are explicit.

For natural arithmetic, the witness is the undefined seed $0/0$ together with the context $x \mapsto x + 0.$ If a finite code of bound m existed, applying the bounded pigeonhole theorem to the $m + 1$ inputs $0, \dots, m$ would give distinct i, j with equal codes, contradicting injectivity. The checked statements are `ResolutionSemantics.PropernessCriteria.onePointProperWithoutOldFixing` and `ResolutionSemantics.PropernessCriteria.natProperWithoutFiniteCoding.` $\square$

Proposition 4.4 (Completed operations). *Every generated operation $f_{R(D)}$ is jointly stage-compatible and therefore extends pointwise to a unique jointly nonexpansive operation*

$$\widehat{f} : \widehat{R(D)} \times \widehat{R(D)} \longrightarrow \widehat{R(D)}.$$

For embedded points,

$$\widehat{f}(\eta x, \eta y) = \eta(f_{R(D)}(x, y)).$$

The completed carrier and operations form a compatible total algebra over $D.$ Its completeness is exposed by `ResolutionSemantics.completionComplete`; the compatible algebra is `ResolutionSemantics.completedAlgebra.`

4.1 Relative universal property

The formalization defines a category-like class of targets consisting of:

- a compatible total algebra over $D;$
- a separated descending filtration;
- nonexpansive generated interpretation;
- jointly nonexpansive primitive operations;
- completeness of the filtered carrier.

For every such target $T,$ generated interpretation extends uniquely from the stage-dense subalgebra, and binary-extension uniqueness shows that the extension preserves all primitive operations. Hence the completion is initial relative to this specified target class.

The adjectives “initial,” “reflective,” and “idempotent” are mathematically correct only with these hypotheses. They are consequences of standard completion arguments and are not presented here as separate novelty claims.

5 Exact Equational Conservativity

Let algebraic terms contain variables, named old-carrier constants, and binary operation symbols. They can be evaluated in $R(D)$ and in $\widehat{R(D)}.$ Embedding commutes with term evaluation by induction. Moreover, every term operation preserves each stage relation because the primitive operations do.

Fix a completed assignment ρ and a stage $n.$ Stage-density allows us to choose, independently for every variable, a generated assignment ρ_n with

$$\rho(v) \equiv_n \eta(\rho_n(v)).$$

Stage compatibility then gives

$$\llbracket s \rrbracket_\rho \equiv_n \eta(\llbracket s \rrbracket_{\rho_n})$$

and similarly for t .

Theorem 5.1 (Exact equation preservation). *For every constant-bearing equation $s = t$,*

$$\left(\forall \rho : \text{Var} \rightarrow R(D), \llbracket s \rrbracket_\rho = \llbracket t \rrbracket_\rho \right)$$

if and only if

$$\left(\forall \widehat{\rho} : \text{Var} \rightarrow \widehat{R(D)}, \llbracket s \rrbracket_{\widehat{\rho}} = \llbracket t \rrbracket_{\widehat{\rho}} \right).$$

Proof. For the forward implication, approximate an arbitrary completed assignment at every stage. The generated equation identifies the two embedded approximants, so the completed term values agree at every stage. Separatedness yields exact equality. The reverse implication follows by evaluating on embedded generated assignments and using injectivity of the embedding. $\square$

The public theorem is `ResolutionSemantics.equationConservative`.

Remark 5.2 (Conditional equations). The theorem concerns universal equations. Arbitrary quasi-equations do not follow from density alone. The formal development transfers conditional laws only under explicit stage-robust or dense-premise hypotheses. No unrestricted quasi-equational conservativity claim is made.

6 Arithmetic Examples

The general construction does not depend on ring or field axioms. We include natural-number and integer instances because they make the behavior at a zero denominator concrete.

6.1 Natural numbers

Take operations $+$, $\cdot$, and $/$ on $\mathbb{N}$, with division undefined exactly when the denominator is zero. For $b \neq 0$,

$$\text{divide}(a, b) = \text{old}(a/b).$$

At zero,

$$\text{divide}(a, 0) = \text{susp}(\text{div}, \text{old}(a), \text{old}(0)).$$

Therefore

$$a/0 = b/0 \implies a = b,$$

and no $a/0$ is an old natural number. These properties remain true after the completion embedding. They are exposed by `ResolutionSemantics.NatDivision.singularFamilyInjective` and `ResolutionSemantics.NatDivision.singularFamilyDisjoint`.

The observational completion is nontrivial in a stronger sense. By [theorem 4.2](#), repeatedly adding the old right identity to $0/0$ produces distinct finite Answers that no fixed observation stage can fully separate. The factorially sampled family is Cauchy without a generated limit, so the completion contains an explicit point outside the generated image, and the explicit pair at budget m has separation rank greater than m . This uses the old-fixing-context criterion rather than the finite-base theorem.

6.2 Integers

The integer instance additionally includes subtraction. Again, every nonzero division agrees with the old evaluator, while

$$a/0 = \text{susp}(\text{div}, \text{old}(a), \text{old}(0)).$$

The completed singular family is injective, and completed $0/0$ is not any embedded integer. The paper-facing results are `ResolutionSemantics.IntDivision.singularFamilyInjective` and `ResolutionSemantics.IntDivision.zeroDivZeroNotOld`.

Integer addition also has old right identity zero. Hence the same old-fixing-context criterion proves that every finite integer-arithmetic stage is strictly coarser than equality, its separation ranks are unbounded, and its completion embedding is not surjective. These corollaries are `ResolutionSemantics.IntDivision.everyFiniteStageNotEquality`, `ResolutionSemantics.IntDivision.factorialOldFixingCauchy`, and `ResolutionSemantics.IntDivision.completionEmbeddingNotSurjective`.

The construction does not assign an ordinary numerical value to $0/0$. It gives $0/0$ a structured Answer in a conservative extension.

This distinction matters when comparing with totalized arithmetic. In an involutive meadow one usually sets $0^{-1} = 0$; in a common meadow exceptional terms propagate one designated error; wheels adopt a different algebraic extension with infinity- and error-like behavior [BHT09, BP21, Car04]. Resolution chooses provenance retention rather than a fixed old value or one undifferentiated error.

7 Relation to Prior Work

7.1 Free completion and partial algebras

The closest structural prior art is the free completion of partial algebras and the Schmidt construction [VL23]. Theorem 2.3 should be read as a Lean-checked concrete presentation of that general idea for the one-sorted binary setting. Other completion and encoding results show that partial-to-total extension is not in itself a novel contribution [Dia09, BKdV96, BKdV99, Nes11].

7.2 Operation trees and effects

Separating operation syntax from interpretation is standard in term algebras, free monads, and algebraic-effects semantics [PP03]. We do not claim the first operation-tree semantics. The specific normalization discipline is: reduce an old-defined node immediately and retain an old-undefined node with its already resolved children.

7.3 Tree metrics and completion

Ultrametric and algebraic treatments of infinite trees long predate this work [AN80, Cou83]. Our completion is distinguished by its source of finite observations: agreement in all

compatible finite-tag models, rather than prefix depth. The Cauchy-sequence quotient and continuous extension arguments are standard once the separated filtration is available.

The distinction is not a matter of presentation. Write $x \equiv_n^d y$ for agreement of two Answers on all levels strictly above depth n ; this is the filtration whose completion is the classical infinite-tree space. The two uniformities differ, and a single family of witnesses shows it.

Proposition 7.1. *Let D be a binary partial algebra with $\text{eval}_D(f, a, z) = \text{none}$ for some f, a, z . Define the combs by $c_0 = a$ and $c_{k+1} = f(c_k, z)$. Then the sequence $(\text{Res}_D c_k)_{k \in \mathbb{N}}$ is Cauchy for the filtration $\equiv_n^d$ but not Cauchy for the filtration $\equiv_n$.*

The formal statement is `Resolution.Probe.observational_strictly_finer_than_depth`. Each comb suspends at every level, so $\text{Res}_D c_i$ and $\text{Res}_D c_j$ agree above depth n whenever $i, j \geq n$: under the tree metric the sequence converges to the infinite comb. For the finite-tag filtration one preserving model with a *single* fresh tag suffices to break it. Send an undefined old pair to the tag, the tag to $\star$, and $\star$ back to the tag; the fold along the combs then alternates between the tag and $\star$, so adjacent combs are separated at stage 1 for every k , however deep the disagreement lies (`Resolution.Probe.comb_separatesAt_one`). A fixed observational budget therefore sees through unbounded depth, which prefix truncation cannot do. Consequently the identity on $R(D)$ does not carry $\equiv_n^d$ -Cauchy sequences to $\equiv_n$ -Cauchy ones, and no uniformly continuous isomorphism of the two completions commutes with the canonical embeddings. Only the statement about the two filtrations is formalized; the consequence for the completions is the standard argument on uniformly continuous maps and is not machine-checked.

In the arithmetic instance of [section 6](#) the effect is sharper: the combs are the iterated singular divisions $(\cdots((a/0)/0)\cdots/0)$, and *no fresh tags at all* are needed. Sending each singular division to the successor of its numerator is a preserving model over $\mathbb{N} \sqcup \{\star\}$ that separates every pair of distinct combs at stage 0, letting the old carrier itself do the counting (`Resolution.Probe.NatProbe.natComb_pairwise_separated_at_zero`).

That stage-zero separator does not make the full natural-arithmetic filtration discrete. It exploits fresh old naturals generated by the iterated-division table. The different family $c_0 = 0/0$, $c_{k+1} = c_k + 0$ cannot use such tags: preservation forces $a + 0 = a$ on every old natural. In fact c_1 and c_2 agree in every stage-zero observer, and [theorem 4.2](#) gives an indistinguishable distinct pair at every finite stage.

The iterated-division comb also bears on the sharpness of the bound in [theorem 3.3](#). Since $|\text{Res}_D c_k| = 2k + 1$, that bound permits $4k + 4$ tags to separate c_k from c_{k+1} , whereas one tag suffices in general and none in arithmetic. The linear budget is therefore far from tight on deep pairs; we do not know its exact order of growth.

7.4 Finite complements and profinite comparison

Finite-complement extensions are classical in semigroup theory under the name of finite Rees index. Cain and Maltcev’s survey explicitly records that residual finiteness passes both to finite-Rees-index subsemigroups and to their finite-complement extensions, citing the theorem of Ruškuc and Thomas [\[CM14, RT98\]](#). Thus residual finiteness is not a gap in that literature. Miller, O’Reilly, Quick, and Ruškuc study residual, weak, strong, and complete separability conditions for semigroups in the finite-quotient tradition [\[MOQR21\]](#).

These results are important adjacent prior art, but neither citation by itself identifies the observer class used here. Our base is a pointwise-fixed partial subalgebra, while a separator

is a compatible total extension whose complement over that base is finite; the separating codomain need not itself be finite when the base is infinite. Moreover, ordinary Rees index starts with a closed subsemigroup. By [proposition 3.6](#), the old image in the generated Answer algebra is closed under the total operations exactly in the already total case. Hence finite Rees index is not literally the inclusion used by a genuinely partial Resolution algebra. The remaining specialist question is whether separation by finite-complement extensions of a pointwise-fixed *partial* base has already been studied under a generalized Rees-index, relative residual-finiteness, or separability formulation.

There is likewise a mature theory of profinite completions and natural extensions of residually finite algebras, as well as categorical formulations through profinite monads [[DaGHP11](#), [CAMU16](#)]. Recent work on profinite trees gives a particularly close comparison on the tree side [[Mor26](#)]. Those constructions are organized by finite quotients or finite recognizing algebras. Our stages instead allow a possibly infinite preserved base and bound only the genuinely new complement. The finite-base properness theorem above is therefore not presented as a theorem about ordinary profinite completion. The documented search has found these analogues, but no exact equivalence or prior statement of the criterion; specialist review is still required.

7.5 Totalized division

Meadows, common meadows, wheels, non-involutive meadows, partial meadows, and fracpair constructions provide mature alternatives with strong equational and initial-algebra results [[BT07](#), [BHT09](#), [BM15](#), [BP21](#), [Car04](#), [BP16](#)]. Resolution is not proposed as uniformly superior. Its design goal is different: preserve old-defined calculations while retaining the provenance of old-undefined applications.

7.6 Candidate contribution

The narrow evidence-qualified contribution is the following combined package:

1. a normalized free compatible completion with structured undefined nodes;
2. an intrinsic observer class whose complement over the preserved base injects into $\text{Fin}(n + 1)$, together with its derived canonical finite-tag presentation;
3. pairwise intrinsic finite-complement separation, with a baseline bound and much smaller witnesses on comb families;
4. the induced filtration and an exact finite-base criterion for when its completion is proper;
5. an infinite-base old-fixing-context criterion proving properness, failure of finite-stage discreteness, and unbounded separation rank, with natural and integer arithmetic as corollaries;
6. a proof that the finite-base and old-fixing-context properness hypotheses are incomparable;
7. completed operations and exact universal equational conservativity;
8. a complete Lean formalization of the package.

Our search has not found this exact combination. This statement reports the current search result, not proof that no earlier equivalent construction exists.

8 Formalization and Trust Boundary

The paper-only repository exposes a deliberately small public interface:

```
import ResolutionSemanticsCompletion
import ResolutionIntrinsicFiniteComplementAPI
import ResolutionFiniteBasePropernessPublic
import ResolutionOldFixingContextPropernessPublic
import ResolutionPropernessCriteriaComparison
import ResolutionArithmeticPropernessPublic

#check ResolutionSemantics.freeCompletionUniversal
#check ResolutionSemantics.OutsideOld
#check ResolutionSemantics.oldOrOutsideEquiv
#check ResolutionSemantics.outsideEquiv
#check ResolutionSemantics.qualitativeFiniteComplementSeparating_theorem
#check ResolutionSemantics.intrinsicExtensionHasCanonicalPresentation
#check ResolutionSemantics.finiteExternalSeparation
#check ResolutionSemantics.finiteSeparationBound
#check ResolutionSemantics.completionComplete
#check ResolutionSemantics.equationConservative
#check ResolutionSemantics.FiniteBaseCompletion.completeIffTotal
#check ResolutionSemantics.FiniteBaseCompletion.
  embeddingNotSurjectiveIffExistsUndefined
#check ResolutionSemantics.OldFixingContextCompletion.embeddingNotSurjective
#check ResolutionSemantics.OldFixingContextCompletion.
  separationRankGreaterThanBudget
#check ResolutionSemantics.PropernessCriteria.onePointProperWithoutOldFixing
#check ResolutionSemantics.PropernessCriteria.natProperWithoutFiniteCoding
#check ResolutionSemantics.NatDivision.separationRanksUnbounded
#check ResolutionSemantics.IntDivision.completionEmbeddingNotSurjective
#check ResolutionSemantics.IntDivision.separationRanksUnbounded
```

The repository pins Lean 4.33.0. The command

```
bash scripts/verify.sh
```

compiles the checked-in proof dependency chain, checks every manuscript-linked declaration, rejects direct proof-hole commands, and prints the axiom dependencies of the principal public theorems. The separate command `bash scripts/build-paper.sh` rebuilds and inspects this PDF. Finite indexing and witness selection use ordinary classical choice in documented places. The result is relative to Lean’s underlying logic and implementation; no foundation-independent consistency claim is made.

The paper gives full mathematical arguments for the free-completion and finite separator constructions because those determine its positioning. Routine completion lemmas are summarized and linked to exact declarations in `THEOREM_MAP.md`.

9 Limitations and Open Questions

1. The formalization is currently one-sorted and binary. General many-sorted and arbitrary-arity versions should be compared directly with the established Schmidt

construction.

2. The intrinsic finite-complement class is deliberately tailored to pointwise preservation of the old partial algebra. Its precise relation to known relative residual-finiteness and syntactic-congruence constructions remains under review.
3. The linear tag bound is a baseline upper bound, not an optimality theorem; its relative state complexity is open.
4. Properness is classified exactly for finite old carriers. The old-fixing-context hypothesis is strictly sufficient and incomparable with finite coding; a full arbitrary-carrier classification remains open.
5. The construction is semantic, not an efficiency claim. Retaining full provenance can produce large finite trees.
6. No ordinary field law is asserted globally on singular Answers unless it is separately proved in the generated algebra.
7. Universal equations transfer exactly; conditional equations require additional hypotheses.
8. Historical priority for the combined package remains unresolved until specialist review is complete.

10 Conclusion

Resolution semantics is best understood as a normalized structured presentation of a free compatible completion of a partial algebra. Its basic conservativity follows from the reduction rule. The substantive additional construction is intrinsic finite-complement separation over the pointwise-preserved old fragment. It separates every generated pair while adding only finitely many states beyond the base; a derived finite-tag normal form supplies the explicit witnesses and induces a separated observational filtration. For finite bases this implies ordinary residual finiteness with an explicit whole-target bound. For genuinely partial algebras the old image is not a total subalgebra, so the finite-complement condition is not literally ordinary Rees index. The general linear estimate is a baseline; its sharper relative complexity is open. Over every finite old carrier, the completion is proper exactly when an old application is undefined. This finite-base hypothesis and the old-fixing-context hypothesis are incomparable, as witnessed by the one-point algebra and natural arithmetic respectively. Standard completion methods then produce a complete total algebra with canonical operations and unchanged universal equational theory.

The old-fixing-context theorem explains the arithmetic examples uniformly: for both natural and integer arithmetic every finite stage is strictly coarser than equality, the separation ranks are unbounded, and the factorial repeated- $+0$ family supplies an explicit completion point outside the generated image. At the same time, $0/0$ remains a structured element outside the old numerical image; the construction does not replace ordinary arithmetic with a new numerical convention. The formal results are reproducible in Lean, while the novelty claim is intentionally limited to the specific verified combination described above.

References

- [AN80] Andre Arnold and Maurice Nivat. The metric space of infinite trees: Algebraic and topological properties. *Fundamenta Informaticae*, 3(4):445–475, 1980.
- [BHT09] Jan A. Bergstra, Yoram Hirshfeld, and John V. Tucker. Meadows and the equational specification of division. *Theoretical Computer Science*, 410(12–13):1261–1271, 2009.
- [BKdV96] Inge Bethke, Jan Willem Klop, and Roel C. de Vrijer. Completing partial combinatory algebras with unique head-normal forms. In *Proceedings of the Eleventh Annual IEEE Symposium on Logic in Computer Science*, pages 448–454. IEEE Computer Society Press, 1996.
- [BKdV99] Inge Bethke, Jan Willem Klop, and Roel C. de Vrijer. Extending partial combinatory algebras. *Mathematical Structures in Computer Science*, 9(4):483–505, 1999.
- [BM15] Jan A. Bergstra and C. A. Middelburg. Division by zero in non-involutive meadows. *Journal of Applied Logic*, 13(1):1–12, 2015.
- [BP16] Jan A. Bergstra and Alban Ponse. Fracpairs and fractions over a reduced commutative ring. *Indagationes Mathematicae*, 27(3):727–748, 2016.
- [BP21] Jan A. Bergstra and Alban Ponse. Division by zero in common meadows, 2021. Version 4.
- [BT07] Jan A. Bergstra and John V. Tucker. The rational numbers as an abstract data type. *Journal of the ACM*, 54(2):7:1–7:25, 2007.
- [CAMU16] Liang-Ting Chen, Jiří Adámek, Stefan Milius, and Henning Urbat. Profinite monads, profinite equations, and reiterman’s theorem. In *Foundations of Software Science and Computation Structures*, volume 9634 of *Lecture Notes in Computer Science*, pages 531–547. Springer, 2016.
- [Car04] Jesper Carlström. Wheels—on division by zero. *Mathematical Structures in Computer Science*, 14(1):143–184, 2004.
- [Cha20] Zachary Chase. A new upper bound for separating words, 2020.
- [CM14] Alan J. Cain and Victor Maltcev. For a few elements more: A survey of finite rees index. *International Journal of Algebra and Computation*, 24(5):611–638, 2014.
- [Cou83] Bruno Courcelle. Fundamental properties of infinite trees. *Theoretical Computer Science*, 25(2):95–169, 1983.
- [DaGHP11] Brian A. Davey, Maria Jo ao Gouveia, Miroslav Haviar, and Hilary A. Priestley. Natural extensions and profinite completions of algebras. *Algebra Universalis*, 66:205–241, 2011.
- [Dia09] Razvan Diaconescu. An encoding of partial algebras as total algebras. *Information Processing Letters*, 109(23–24):1245–1251, 2009.

- [GK86] Pavel Goralcik and Vaclav Koubek. On discerning words by automata. In *Automata, Languages and Programming: 13th International Colloquium, ICALP 1986*, volume 226 of *Lecture Notes in Computer Science*, pages 116–122. Springer, 1986.
- [MOQR21] Craig Miller, Gerard O’Reilly, Martyn Quick, and Nik Ruškuc. On separability finiteness conditions in semigroups, 2021.
- [Mor26] Vincent Moreau. Profinite trees, through lawvere theories and the lambda-calculus, 2026. Version 2.
- [Nes11] Gabriela Asli Rino Nesin. Completing partial algebra models of term rewriting systems. Master of logic thesis, Institute for Logic, Language and Computation, University of Amsterdam, 2011.
- [PP03] Gordon D. Plotkin and John Power. Algebraic operations and generic effects. *Applied Categorical Structures*, 11(1):69–94, 2003.
- [Rob89] J. M. Robson. Separating strings with small automata. *Information Processing Letters*, 30(4):209–214, 1989.
- [RT98] Nik Ruškuc and Richard M. Thomas. Syntactic and Rees indices of subsemi-groups. *Journal of Algebra*, 205(2):435–450, 1998.
- [VL23] Juan Climent Vidal and Enric Cosme Llopez. Functoriality of the schmidt construction. *Logic Journal of the IGPL*, 31(5):822–893, 2023.